

WEEK 3 TERADATA EXERCISE GUIDE

Specific Goals for Week 3 Teradata Exercises

Use these exercises to:

- Understand the different subsets of sku numbers and departments in each of the tables in the Dillard's database (it is important to understand this in order to draw conclusions about Dillard's sales patterns)
- Learn the differences between Teradata's and MySQL's GROUP BY syntax
- Practice joining tables using a relational schema as a guide

GROUP BY clauses in Teradata vs. MySQL

The syntax for standard aggregate functions, HAVING clauses, and joins is the same in Teradata as it is in MySQL. However, GROUP BY clauses in Teradata differ from GROUP BY clauses in MySQL. In MySQL Exercises 5 and 6, we discussed how the outputs of a query must all be aggregated at the same level. If you include a GROUP BY clause in your query, but forget to aggregate a column included in your SELECT list, MySQL will output a random row from the un-aggregated column when the query is executed. This functionality is convenient when you know that all the values within a group of an un-aggregated column are the same, but also makes it easy to misinterpret the results of poorly written queries. To prevent users from having the opportunity to write poor queries in the first place, Teradata has a different rule for GROUP BY statements than does MySQL. In Teradata (and many other database systems):

Any non-aggregate column in the SELECT list or HAVING list of a query with a GROUP BY clause must also listed in the GROUP BY clause.

This rule makes it less convenient to write some queries, but ensures that outputs that aggregate across groups of records aren't mixed with outputs that only represent single records.

To understand the aggregation rule better, consider this query:

```
SELECT sku, COUNT(sku), retail, cost
FROM skstinfo
GROUP BY sku
```

This query would be executed in MySQL (if the Dillard's dataset was a MySQL database), but a random value in the retail column and the cost column would be outputted for each output row that reports a count of a given sku number. Teradata, on the other hand, will return this error if you tried to execute the query above:

Error Code - 3504

Error Message - [Teradata Database] [TeraJDBC 15.10.00.05] [Error 3504] [SQLState HY000] Selected non-aggregate values must be part of the associated group.

To resolve the error, either retail and cost both need to be included in the GROUP BY clause (but note that this query would output a separate row for each distinct combination of sku, retail, and cost rather than a row for each sku, on its own):

```
SELECT sku, retail, cost, COUNT(sku)
FROM skstinfo
GROUP BY sku, retail, cost
```

...or retail and cost need to be aggregated to match the sku aggregation:

```
SELECT sku, COUNT(sku), AVG(retail), AVG(cost)
FROM skstinfo
GROUP BY sku
```

The “Selected non-aggregate values must be part of the associated group” error is one of the most common errors new SQL learners make, so keep an eye out for it. If you see it, immediately check the columns you have included in your SELECT and GROUP BY clauses.

Integrate everything you learned this week

Keep your Dillard’s relational schema handy while you work through these exercises. You will likely need to look up column names, and will want to refer to the diagram to determine which columns to use to join your tables.

Exercise 1: (a) Use COUNT and DISTINCT to determine how many distinct skus there are in pairs of the skuinfo, skstinfo, and trnsact tables. Which skus are common to pairs of tables, or unique to specific tables?

The significance of your answers will become clear in Exercise 3.

You will only be tested on queries that compare 2 skus in 2 of the tables listed above at a time. However, if you want to try writing queries that compare all 3 tables at once, you might find the following references helpful:

<http://www.w3resource.com/sql/joins/perform-a-left-join.php>

<http://www.wellho.net/solutions/mysql-left-joins-to-link-three-or-more-tables.html>

After interpreting the set of queries you decide to implement, you should conclude that only one of the three tables has distinct skus that are not contained in either of the other tables. However, none of the three tables contain the exact same number of distinct skus.

Note that any queries that join on the `trnsact` table will likely take a while to run.

(b) Use COUNT to determine how many instances there are of each sku associated with each store in the `skstinfo` table and the `trnsact` table?

Note that these queries will take a while to run.

You should see there are multiple instances of every sku/store combination in the `trnsact` table, but only one instance of every sku/store combination in the `skstinfo` table. Therefore you could join the `trnsact` and `skstinfo` tables, but you would need to join them on both of the following conditions: `trnsact.sku= skstinfo.sku AND trnsact.store= skstinfo.store`.

Exercise 2: (a) Use COUNT and DISTINCT to determine how many distinct stores there are in the `strinfo`, `store_msa`, `skstinfo`, and `trnsact` tables.

You should see that:

distinct stores in `strinfo` > # distinct stores in `skstinfo` > # distinct stores in `store_msa` > # distinct stores in `trnsact`

Use this information to help you assess which table to query when answering questions about Dillard's stores in the quiz.

(b) Which stores are common to all four tables, or unique to specific tables?

Refer to the links provided in Exercise 1a to practice writing queries that compare all 4 tables at once.

Exercise 3: It turns out there are many skus in the `trnsact` table that are not in the `skstinfo` table. As a consequence, we will not be able to complete many desirable analyses of Dillard's profit, as opposed to revenue, because we do not have the cost information for all the skus in the `trnsact` table (recall that $\text{profit} = \text{revenue} - \text{cost}$). Examine some of the rows in the `trnsact` table that are not in the `skstinfo` table; can you find any common features that could explain why the cost information is missing?

Please note: The join you will need to complete this analysis will take a long time to run. This query will give you a good feeling for what working with enterprise-sized data feels like. You might want to pin the results once you retrieve them so that you can examine the results later in the session.

Exercise 4: Although we can't complete all the analyses we'd like to on Dillard's profit, we can look at general trends. What is Dillard's average profit per day?

We can justify examining the *average* profit per day (excluding transactions that cannot be joined with the skstinfo), but we cannot justify examining the *total* profit over a given timeframe because there are too many transactions for which we do not have cost information.

To calculate profit, subtract the total cost of the items in a transaction from the total revenue brought in during the transaction. Refer to the Dillard's Department Store Database Metadata and what you've seen so far to determine what column—or columns—represents total revenue, and what columns are needed to calculate the total cost associated with a transaction (note that some transactions have multiple items, so you will need to take that into account in your calculations, especially your cost calculations). If you are interested in looking at the total value of goods purchased or returned, use the “amt” field. If you are interested in looking at the total number of goods purchased or returned, use the “quantity” field. Make sure to specify the correct stype in your query. Pay close attention to what we learned in Exercise 1b.

If you are calculating the average profit per day correctly, you will find that the average profit per day from register 640 is \$10,779.20.

Exercise 5: On what day was the total value (in \$) of returned goods the greatest? On what day was the total number of individual returned items the greatest?

Make sure to specify the correct stype in your query. If you are interested in looking at the total value of goods purchased or returned, use the “amt” field. If you are interested in looking at the total number of goods purchased or returned, use the “quantity” field.

Exercise 6: What is the maximum price paid for an item in our database? What is the minimum price paid for an item in our database?

Make sure to use the correct column to address these questions (use the same column(s) as you used to calculate revenue). The answer to the minimum price question is likely evidence of more incorrect entries.

Exercise 7: How many departments have more than 100 brands associated with them, and what are their descriptions?

A HAVING clause will be helpful for addressing this question. You will also need a join to combine the skuinfo and deptinfo tables in order to retrieve the descriptions of the departments.

Exercise 8: Write a query that retrieves the department descriptions of each of the skus in the skstinfo table.

You will need to join 3 tables in this query. Start by writing a query that connects the skstinfo and skuinfo tables. Once you are sure that query works, connect the deptinfo table. You may want to explore what happens

when you incorporate aggregate functions into your query. When you do so, make sure all of your column names are referenced by the correct table aliases.

If you have written your query correctly, you will find that the department description for sku #5020024 is “LESLIE”.

Exercise 9: What department (with department description), brand, style, and color had the greatest total value of returned items?

You will need to join 3 tables in this query. Start by writing a query that connects 2 tables. Once you are sure that query works, connect the 3rd table. Make sure you include both of these fields in the SELECT and GROUP BY clauses. Make sure any non-aggregate column in your SELECT list is also listed in your GROUP BY clause.

If you have written your query correctly, you will find that the department with the 5th highest total value of returned items is #2200 with a department description of “CELEBRT”, a brand of “LANCOME”, a style of “1042”, a color of “00-NONE”, and a total value of returned items of \$177,142.50.

Exercise 10: In what state and zip code is the store that had the greatest total revenue during the time period monitored in our dataset?

You will need to join two tables to answer this question, and will need to divide your answers up according to the “state” and “zip” fields. Make sure you include both of these fields in the SELECT and GROUP BY clauses. Don’t forget to specify that you only want to examine purchase transactions (not returns).

If you have written your query correctly, you will find that the department with the 10th highest total revenue is in Hurst, TX.

What to do when you don’t know how to answer an exercise

If you are having trouble writing some of your queries, don’t worry! Here are some things you can try:

1. Make sure your query is consistent with the items listed in the Syntax Error Checklist. You can copy and paste the URL below or look in course resources for the checklist.

<https://www.coursera.org/learn/analytics-mysql/resources/AaIGu>

2. Break down your query into the smallest pieces, or clauses, possible. Once you make sure each small piece works on its own, add in another piece or clause, one by one, always making sure the modified query works before you add in another clause.

3. Search the previous discussion forum posts to see if other students had similar questions or challenges.
4. If you've exhausted all your other options, ask for help in the Discussion forums. Remember to (a) list the exercise name and question number in the title of the post or at the very beginning of the post, and (b) copy and paste the text of the question you are trying to answer at the beginning of the post. If you would like help troubleshooting a query, include what query (or queries) you tried, the error(s) you received, and the logic behind why you wrote the query the way you did.

Test your understanding using this week's graded quiz once you feel you fully understand how to do the following in both MySQL and Teradata:

- Summarize your data using the aggregate functions MIN, MAX, SUM, and AVG
- Count the number of total rows, and the number of distinct rows, in your data
- Group your data and/or summaries according to values in a column
- Restrict the groups you retrieve according to specific criteria
- Combine information from two tables using inner, left, and right joins
- Combine information from more than two tables using inner, left, and right joins